

# Shangbin Feng

Paul G. Allen School of Computer Science & Engineering  
University of Washington  
185 E Stevens Way NE, Seattle, WA 98195

Homepage: [bunsenfeng.github.io](https://bunsenfeng.github.io)  
Email: [shangbin@cs.washington.edu](mailto:shangbin@cs.washington.edu)  
Scholar: [Google Scholar](#)

## Education

---

|                 |                                                                                                                                                                                                                                                                           |
|-----------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| 09/2022–Present | Ph.D. candidate, Computer Science & Engineering, University of Washington<br>Advisor: Yulia Tsvetkov<br>Thesis: <i>Model Collaboration, For A Bottom-Up and Participatory AI Future</i><br>Committee: Yulia Tsvetkov, Yejin Choi, Luke Zettlemoyer, Colin Raffel, Fei Xia |
| 09/2022–03/2024 | M.S., Computer Science & Engineering, University of Washington                                                                                                                                                                                                            |
| 09/2018–07/2022 | B.E., Computer Science and Technology, Xi'an Jiaotong University                                                                                                                                                                                                          |

## Honors & Awards

---

|           |                                                                        |
|-----------|------------------------------------------------------------------------|
| 2026      | Senior Area Chair Award, ACL 2026                                      |
| 2026      | NVIDIA Graduate Fellowship                                             |
| 2025      | Jane Street Graduate Research Fellowship                               |
| 2024      | IBM PhD Fellowship                                                     |
| 2024      | Outstanding Paper Award, ACL 2024                                      |
| 2024      | Senior Area Chair Award, ACL 2024                                      |
| 2023      | Best Paper Award, ACL 2023                                             |
| 2022–2024 | Top / Outstanding Reviewer (NeurIPS 2022 & 2023, LoG 2022, EMNLP 2024) |
| 2022      | Outstanding Undergraduate Thesis Award                                 |

## Employment

---

|                 |                                                                                      |
|-----------------|--------------------------------------------------------------------------------------|
| 06/2026–09/2026 | Research Intern, NVIDIA, Seattle, WA<br>Managers: David Acuna, Jan Kautz             |
| 06/2024–09/2024 | Student Researcher, Google, Los Angeles, CA<br>Managers: Zifeng Wang, Chen-Yu Lee    |
| 09/2022–12/2024 | Research Assistant, University of Washington, Seattle, WA<br>Manager: Yulia Tsvetkov |

## Publications

---

\* denotes equal contribution. Full list on [Google Scholar](#).

*Under Review*

- [U7] **Shangbin Feng**, Yike Wang, Weijia Shi, Luke Zettlemoyer, Yejin Choi, and Yulia Tsvetkov. “Scaling Participation in Modular AI Systems.” Under review at a journal. 2026.
- [U6] **Shangbin Feng**<sup>\*</sup>, Yuyang Bai<sup>\*</sup>, Ziyuan Yang<sup>\*</sup>, Yike Wang, Zhaoxuan Tan, Jiajie Yan, Zhenyu Lei, Wenxuan Ding, Weijia Shi, Haojin Wang, Zhenting Qi, Yuru Jiang, Heng Wang, Chengsong Huang, Yu Fei, Jihan Yao, Yilun Du, Luke Zettlemoyer, Yejin Choi, and Yulia Tsvetkov. “MoCo: A One-Stop Shop for Model Collaboration Research.” Under review. 2026.

- [U5] **Shangbin Feng**, Yulia Tsvetkov, Jan Kautz, and David Acuna. “Compositional RL: Evaluating and Training LLMs for Collaborative Reasoning.” Under review. 2026.
- [U4] Christina Hahn\*, **Shangbin Feng**\*, Dean Light, Swastik Roy, Hila Gonen, and Yulia Tsvetkov. “Multi-LLM Collaborative Alignment via Stackelberg Games.” Under review. 2026.
- [U3] Ziyuan Yang\*, Yike Wang\*, **Shangbin Feng**\*, and Yulia Tsvetkov. “Save Your Saturated Data: Learning Beyond Reward Saturation in Group-Based RL.” Under review. 2026.
- [U2] Zongwan Cao\*, Ziyuan Yang\*, **Shangbin Feng**\*, Michael Duan, Skyler Hallinan, Bingbing Wen, Lucy Lu Wang, and Yulia Tsvetkov. “You’re Hired: Strategic Model Selection for LLM Collaboration.” Under review. 2026.
- [U1] Jihan Yao, Sihan Zeng, **Shangbin Feng**, Zhiyuan Fan, Banghua Zhu, and Yulia Tsvetkov. “Before the Rollout Ends: Early Terminal Reward Prediction for Long-horizon Coding Agents.” Under review. 2026.

### *Conference and Journal Publications*

- [C53] Zhuofeng Li, Haoxiang Zhang, Cong Wei, Pan Lu, Ping Nie, Yi Lu, Yuyang Bai, **Shangbin Feng**, Hangxiao Zhu, Ming Zhong, Yuyu Zhang, Jianwen Xie, Yejin Choi, James Zou, Jiawei Han, Wenhui Chen, Jimmy Lin, Dongfu Jiang, and Yu Zhang. “Beyond Semantic Similarity: Rethinking Retrieval for Agentic Search via Direct Corpus Interaction.” In: *Proceedings of NeurIPS*. 2026. **Spotlight**
- [C52] Yike Wang, Faeze Brahman, **Shangbin Feng**, Teng Xiao, Hannaneh Hajishirzi, and Yulia Tsvetkov. “Small Reward Models via Backward Inference.” In: *Proceedings of EMNLP*. 2026.
- [C51] Zhenyu Lei, Zihan Chen, Yaochen Zhu, **Shangbin Feng**, Zaiyi Zheng, Ruocheng Guo, Yushun Dong, and Jundong Li. “TeacherGRPO: Closing the Capacity Gap in Reasoning Distillation via Teacher Alignment.” In: *Findings of EMNLP*. 2026.
- [C50] Yike Wang, **Shangbin Feng**, Yulia Tsvetkov, and Hannaneh Hajishirzi. “ScienceMeter: Tracking Scientific Knowledge Updates in Language Models.” In: *Proceedings of COLM*. 2026.
- [C49] Jihan Yao, Yushi Hu, Wenyuan Wang, Bin Han, **Shangbin Feng**, Guang Yang, Yujie Yi, Bingbing Wen, Ranjay Krishna, Lucy Lu Wang, Yulia Tsvetkov, Noah A. Smith, and Banghua Zhu. “MMMGen: A Comprehensive and Reliable Benchmark for Multitask Multimodal Generation.” In: *Proceedings of COLM*. 2026.
- [C48] Ziyuan Yang\*, Wenxuan Ding\*, **Shangbin Feng**\*, and Yulia Tsvetkov. “Among Us: Measuring and Mitigating Malicious Contributions in Model Collaboration Systems.” In: *Proceedings of ACL*. 2026. **Senior Area Chair Award; Oral**
- [C47] **Shangbin Feng**, Wenxuan Ding, Alisa Liu, Zifeng Wang, Weijia Shi, Yike Wang, Zejiang Shen, Xiaochuang Han, Hunter Lang, Chen-Yu Lee, Tomas Pfister, Yejin Choi, and Yulia Tsvetkov. “When One LLM Drools, Multi-LLM Collaboration Rules.” In: *Proceedings of ACL*. 2026. **Oral**
- [C46] **Shangbin Feng**, Yike Wang, Weijia Shi, and Yulia Tsvetkov. “Data Swarms: Optimizable Generation of Synthetic Evaluation Data.” In: *Findings of ACL*. 2026.
- [C45] Yizhuo Zhang\*, Heng Wang\*, **Shangbin Feng**\*, Zhaoxuan Tan, Xinyun Liu, and Yulia Tsvetkov. “Generalizable LLM Learning of Graph Synthetic Data with Post-training Alignment.” In: *Findings of ACL*. 2026.
- [C44] Rock Yuren Pang, K. J. Kevin Feng, **Shangbin Feng**, Chu Li, Weijia Shi, Yulia Tsvetkov, Jeffrey Heer, and Katharina Reinecke. “Interactive Reasoning: Visualizing and Controlling Chain-of-Thought Reasoning in Large Language Models.” In: *Proceedings of IUI*. 2026.
- [C43] Yuru Jiang\*, Wenxuan Ding\*, **Shangbin Feng**\*, Greg Durrett, and Yulia Tsvetkov. “Sparta Alignment: Collectively Aligning Multiple Language Models through Combat.” In: *Proceedings of*

- [C42] **Shangbin Feng**, Zifeng Wang, Palash Goyal, Yike Wang, Weijia Shi, Huang Xia, Hamid Palangi, Luke Zettlemoyer, Yulia Tsvetkov, Chen-Yu Lee, and Tomas Pfister. “Heterogeneous Swarms: Jointly Optimizing Model Roles and Weights for Multi-LLM Systems.” In: *Proceedings of NeurIPS*. 2025.
- [C41] Yiwei Yang, Chung Peng Lee, **Shangbin Feng**, Dora Zhao, Bingbing Wen, Anthony Z. Liu, Yulia Tsvetkov, and Bill Howe. “Escaping the SpuriVerse: Can Large Vision-Language Models Generalize Beyond Seen Spurious Correlations?” In: *Proceedings of NeurIPS (Datasets and Benchmarks Track)*. 2025.
- [C40] **Shangbin Feng**, Zifeng Wang, Yike Wang, Sayna Ebrahimi, Hamid Palangi, Lesly Miculicich, Achin Kulshrestha, Nathalie Rauschmayr, Yejin Choi, Yulia Tsvetkov, Chen-Yu Lee, and Tomas Pfister. “Model Swarms: Collaborative Search to Adapt LLM Experts via Swarm Intelligence.” In: *Proceedings of ICML*. 2025.
- [C39] Jillian Fisher, Ruth E. Appel, Chan Young Park, Yujin Potter, Liwei Jiang, Taylor Sorensen, **Shangbin Feng**, Yulia Tsvetkov, Margaret E. Roberts, Jennifer Pan, Dawn Song, and Yejin Choi. “Position: Political Neutrality in AI Is Impossible—But Here Is How to Approximate It.” In: *Proceedings of ICML (Position Track)*. 2025. **Oral**
- [C38] Jillian Fisher, **Shangbin Feng**, Robert Aron, Thomas Richardson, Yejin Choi, Daniel W. Fisher, Jennifer Pan, Yulia Tsvetkov, and Katharina Reinecke. “Biased LLMs can Influence Political Decision-Making.” In: *Proceedings of ACL*. 2025. **Oral**
- [C37] Qingkai Zeng, Yuyang Bai, Zhaoxuan Tan, Zhenyu Wu, **Shangbin Feng**, and Meng Jiang. “Code-Taxo: Enhancing Taxonomy Expansion with Limited Examples via Code Language Prompts.” In: *Findings of ACL*. 2025.
- [C36] Varich Boonsanong, Vidhisha Balachandran, Xiaochuang Han, **Shangbin Feng**, Lucy Lu Wang, and Yulia Tsvetkov. “FACTS&EVIDENCE: An Interactive Tool for Transparent Fine-Grained Factual Verification of Machine-Generated Text.” In: *Proceedings of NAACL (Demo Track)*. 2025.
- [C35] Jihan Yao\*, Wenxuan Ding\*, **Shangbin Feng\***, Lucy Lu Wang, and Yulia Tsvetkov. “Varying Shades of Wrong: Aligning LLMs with Wrong Answers Only.” In: *Proceedings of ICLR*. 2025.
- [C34] Bingbing Wen, Jihan Yao, **Shangbin Feng**, Chenjun Xu, Yulia Tsvetkov, Bill Howe, and Lucy Lu Wang. “Know Your Limits: A Survey of Abstention in Large Language Models.” In: *Transactions of the Association for Computational Linguistics (TACL)*. 2024.
- [C33] Shuyue Stella Li, Vidhisha Balachandran, **Shangbin Feng**, Jonathan Ilgen, Emma Pierson, Pang Wei Koh, and Yulia Tsvetkov. “MediQ: Question-Asking LLMs and a Benchmark for Reliable Interactive Clinical Reasoning.” In: *Proceedings of NeurIPS*. 2024.
- [C32] **Shangbin Feng**, Weijia Shi, Yike Wang, Wenxuan Ding, Orevaoghene Ahia, Shuyue Stella Li, Vidhisha Balachandran, Sunayana Sitaram, and Yulia Tsvetkov. “Teaching LLMs to Abstain across Languages via Multilingual Feedback.” In: *Proceedings of EMNLP*. 2024.
- [C31] **Shangbin Feng**, Taylor Sorensen, Yuhao Liu, Jillian Fisher, Chan Young Park, Yejin Choi, and Yulia Tsvetkov. “Modular Pluralism: Pluralistic Alignment via Multi-LLM Collaboration.” In: *Proceedings of EMNLP*. 2024.
- [C30] Yizhuo Zhang\*, Heng Wang\*, **Shangbin Feng\***, Zhaoxuan Tan, Xiaochuang Han, Tianxing He, and Yulia Tsvetkov. “Can LLM Graph Reasoning Generalize beyond Pattern Memorization?” In: *Findings of EMNLP*. 2024.
- [C29] Qingkai Zeng\*, Yuyang Bai\*, Zhaoxuan Tan, **Shangbin Feng**, Zhenwen Liang, Zhihan Zhang, and Meng Jiang. “Chain-of-Layer: Iteratively Prompting Large Language Models for Taxonomy Induction from Limited Examples.” In: *Proceedings of CIKM*. 2024.

- [C28] Yike Wang\*, **Shangbin Feng\***, Heng Wang, Weijia Shi, Vidhisha Balachandran, Tianxing He, and Yulia Tsvetkov. “Resolving Knowledge Conflicts in Large Language Models.” In: *Proceedings of COLM*. 2024.
- [C27] **Shangbin Feng**, Weijia Shi, Yike Wang, Wenxuan Ding, Vidhisha Balachandran, and Yulia Tsvetkov. “Don’t Hallucinate, Abstain: Identifying LLM Knowledge Gaps via Multi-LLM Collaboration.” In: *Proceedings of ACL*. 2024. **Outstanding Paper Award; Senior Area Chair Award**
- [C26] **Shangbin Feng**, Herun Wan, Ningnan Wang, Zhaoxuan Tan, Minnan Luo, and Yulia Tsvetkov. “What Does the Bot Say? Opportunities and Risks of Large Language Models in Social Media Bot Detection.” In: *Proceedings of ACL*. 2024.
- [C25] Yichen Wang, **Shangbin Feng**, Abe Bohan Hou, Xiao Pu, Chao Shen, Xiaoming Liu, Yulia Tsvetkov, and Tianxing He. “Stumbling Blocks: Stress Testing the Robustness of Machine-Generated Text Detectors Under Attacks.” In: *Proceedings of ACL*. 2024.
- [C24] Wenxuan Ding\*, **Shangbin Feng\***, Yuhan Liu, Zhaoxuan Tan, Vidhisha Balachandran, Tianxing He, and Yulia Tsvetkov. “Knowledge Crosswords: Geometric Knowledge Reasoning with Large Language Models.” In: *Findings of ACL*. 2024.
- [C23] Herun Wan\*, **Shangbin Feng\***, Zhaoxuan Tan, Heng Wang, Yulia Tsvetkov, and Minnan Luo. “DELL: Generating Reactions and Explanations for LLM-Based Misinformation Detection.” In: *Findings of ACL*. 2024.
- [C22] Yuhan Liu\*, **Shangbin Feng\***, Xiaochuang Han, Vidhisha Balachandran, Chan Young Park, Sachin Kumar, and Yulia Tsvetkov. “P<sup>3</sup>SUM: Preserving Author’s Perspective in News Summarization with Diffusion Language Models.” In: *Proceedings of NAACL*. 2024.
- [C21] Yuyang Bai\*, **Shangbin Feng\***, Vidhisha Balachandran, Zhaoxuan Tan, Shiqi Lou, Tianxing He, and Yulia Tsvetkov. “KGQuiz: Evaluating the Generalization of Encoded Knowledge in Large Language Models.” In: *Proceedings of The Web Conference*. 2024. **Oral**
- [C20] **Shangbin Feng**, Weijia Shi, Yuyang Bai, Vidhisha Balachandran, Tianxing He, and Yulia Tsvetkov. “Knowledge Card: Filling LLMs’ Knowledge Gaps with Plug-in Specialized Language Models.” In: *Proceedings of ICLR*. 2024. **Oral (top 1.2%)**
- [C19] **Shangbin Feng**, Vidhisha Balachandran, Yuyang Bai, and Yulia Tsvetkov. “FactKB: Generalizable Factuality Evaluation using Language Models Enhanced with Factual Knowledge.” In: *Proceedings of EMNLP*. 2023.
- [C18] Zhaoxuan Tan\*, **Shangbin Feng\***, Melanie Sclar, Herun Wan, Minnan Luo, Yejin Choi, and Yulia Tsvetkov. “BotPercent: Estimating Bot Populations in Twitter Communities.” In: *Findings of EMNLP*. 2023.
- [C17] Heng Wang, Wenqian Zhang, Yuyang Bai, Zhaoxuan Tan, **Shangbin Feng**, Qinghua Zheng, and Minnan Luo. “Detecting Spoilers in Movie Reviews with External Movie Knowledge and User Networks.” In: *Proceedings of EMNLP*. 2023.
- [C16] Heng Wang\*, **Shangbin Feng\***, Tianxing He, Zhaoxuan Tan, Xiaochuang Han, and Yulia Tsvetkov. “Can Language Models Solve Graph Problems in Natural Language?” In: *Proceedings of NeurIPS*. 2023. **Spotlight (top 3.4%)**
- [C15] **Shangbin Feng**, Chan Young Park, Yuhan Liu, and Yulia Tsvetkov. “From Pretraining Data to Language Models to Downstream Tasks: Tracking the Trails of Political Biases Leading to Unfair NLP Models.” In: *Proceedings of ACL*. 2023. **Best Paper Award** [[Washington Post](#)] [[MIT Technology Review](#)]
- [C14] **Shangbin Feng**, Zhaoxuan Tan, Wenqian Zhang, Zhenyu Lei, and Yulia Tsvetkov. “KALM: Knowledge-Aware Integration of Local, Document, and Global Contexts for Long Document Under-

standing.” In: *Proceedings of ACL*. 2023.

- [C13] Zhenyu Lei\*, Herun Wan\*, Wenqian Zhang, **Shangbin Feng**, Zilong Chen, Jundong Li, Qinghua Zheng, and Minnan Luo. “BIC: Twitter Bot Detection with Text-Graph Interaction and Semantic Consistency.” In: *Proceedings of ACL*. 2023.
- [C12] Yuhao Liu, Zhaoxuan Tan, Heng Wang, **Shangbin Feng**, Qinghua Zheng, and Minnan Luo. “Bot-MoE: Twitter Bot Detection with Community-Aware Mixtures of Modal-Specific Experts.” In: *Proceedings of SIGIR*. 2023.
- [C11] Shujie Yang, Binchi Zhang, **Shangbin Feng**, Zhaoxuan Tan, Qinghua Zheng, Jun Zhou, and Minnan Luo. “AHEAD: A Triple Attention Based Heterogeneous Graph Anomaly Detection Approach.” In: *Proceedings of CIAC*. 2023.
- [C10] Zhaoxuan Tan, Zilong Chen, **Shangbin Feng**, Qingyue Zhang, Qinghua Zheng, Jundong Li, and Minnan Luo. “KRACL: Contrastive Learning with Graph Context Modeling for Sparse Knowledge Graph Completion.” In: *Proceedings of The Web Conference*. 2023.
- [C9] **Shangbin Feng**, Zhaoxuan Tan, Zilong Chen, Peisheng Yu, Ningnan Wang, Qinghua Zheng, Xiaojun Chang, and Minnan Luo. “PAR: Political Actor Representation Learning with Social Context and Expert Knowledge.” In: *Proceedings of EMNLP*. 2022.
- [C8] **Shangbin Feng**\*, Zhaoxuan Tan\*, Herun Wan\*, Ningnan Wang\*, Zilong Chen\*, Binchi Zhang\*, Qinghua Zheng, Wenqian Zhang, Zhenyu Lei, Shujie Yang, Xinshun Feng, Qingyue Zhang, Hongrui Wang, Yuhao Liu, Yuyang Bai, Heng Wang, Zijian Cai, Yanbo Wang, Lijing Zheng, Zihan Ma, Jundong Li, and Minnan Luo. “TwiBot-22: Towards Graph-Based Twitter Bot Detection.” In: *Proceedings of NeurIPS (Datasets and Benchmarks Track)*. 2022.
- [C7] Xinshun Feng, Herun Wan, **Shangbin Feng**, Hongrui Wang, Qinghua Zheng, Jun Zhou, and Minnan Luo. “GraTo: Graph Neural Network Framework Tackling Over-Smoothing with Neural Architecture Search.” In: *Proceedings of CIKM*. 2022.
- [C6] Claudia Flores-Saviaga, **Shangbin Feng**, and Saiph Savage. “Datavoidant: An AI System for Addressing Political Data Voids on Social Media.” In: *Proceedings of CSCW*. 2022.
- [C5] Wenqian Zhang\*, **Shangbin Feng**\*, Zilong Chen\*, Zhenyu Lei, Jundong Li, and Minnan Luo. “KCD: Knowledge Walks and Textual Cues Enhanced Political Perspective Detection in News Media.” In: *Proceedings of NAACL*. 2022. **Oral**
- [C4] **Shangbin Feng**, Zhaoxuan Tan, Rui Li, and Minnan Luo. “Heterogeneity-aware Twitter Bot Detection with Relational Graph Transformers.” In: *Proceedings of AAAI*. 2022.
- [C3] **Shangbin Feng**, Herun Wan, Ningnan Wang, and Minnan Luo. “BotRGCN: Twitter Bot Detection with Relational Graph Convolutional Networks.” In: *Proceedings of ASONAM (Short Paper)*. 2021.
- [C2] **Shangbin Feng**, Herun Wan, Ningnan Wang, Jundong Li, and Minnan Luo. “TwiBot-20: A Comprehensive Twitter Bot Detection Benchmark.” In: *Proceedings of CIKM (Resource Track)*. 2021.
- [C1] **Shangbin Feng**, Herun Wan, Ningnan Wang, Jundong Li, and Minnan Luo. “SATAR: A Self-supervised Approach to Twitter Account Representation Learning and its Application in Bot Detection.” In: *Proceedings of CIKM (Applied Track)*. 2021.

### Other Preprints

- [P3] Haojin Wang, Yike Wang, **Shangbin Feng**, Hannaneh Hajishirzi, and Yulia Tsvetkov. “MentorCollab: Large-to-Small Inference-Time Mentorship for Concise Reasoning in Language Models.” *arXiv preprint*. 2026.
- [P2] Zifeng Zhu\*, **Shangbin Feng**\*, Herun Wan, Ningnan Wang, Minnan Luo, and Yulia Tsvetkov.

“GuessBench: Sensemaking Multimodal Creativity in the Wild.” *arXiv preprint*. 2025.

[P1] Bingbing Wen, Faeze Brahman, Zhan Su, **Shangbin Feng**, Yulia Tsvetkov, Lucy Lu Wang, and Bill Howe. “MARVEL: Modular Abstention for Reliable and Versatile Expert LLMs.” *arXiv preprint*. 2025.

## Invited Talks

---

### “Protocols of Model Collaboration”

|         |                                    |
|---------|------------------------------------|
| 11/2026 | USC (Host: Jonathan May)           |
| 02/2026 | Stanford NLP Seminar               |
| 02/2026 | Northwestern (Host: Manling Li)    |
| 02/2026 | UIUC (Host: Heng Ji)               |
| 02/2026 | Notre Dame (Host: Meng Jiang)      |
| 11/2025 | Columbia (Host: Kathleen McKeown)  |
| 11/2025 | Cornell Tech (Host: Yoav Artzi)    |
| 11/2025 | Northeastern (Host: Weiyan Shi)    |
| 11/2025 | MIT (Host: Yoon Kim)               |
| 11/2025 | Harvard (Host: Yilun Du)           |
| 11/2025 | UPenn (Host: Chris Callison-Burch) |
| 10/2025 | Sahara AI                          |
| 10/2025 | UCLA (Host: Nanyun (Violet) Peng)  |
| 10/2025 | USC (Host: Xiang Ren)              |

### “Model Swarms: Collaborative Search to Adapt LLM Experts via Swarm Intelligence”

|         |                                                   |
|---------|---------------------------------------------------|
| 10/2025 | UW CSE Affiliates Day                             |
| 05/2025 | Jane Street Graduate Research Fellowship Workshop |
| 04/2025 | Google DeepMind                                   |
| 01/2025 | Coding Aperitivo Seminar, Bocconi University      |

### “From Pretraining Data to Language Models to Downstream Tasks: Tracking the Trails of Political Biases Leading to Unfair NLP Models”

|         |                                           |
|---------|-------------------------------------------|
| 11/2023 | UW CSE Affiliates Day                     |
| 07/2023 | ACL 2023 Best Paper Award Plenary Session |
| 06/2023 | NYU Disinformation Symposium              |

## Teaching and Mentoring

---

### *Head Teaching Assistant*

|             |                                                                                                                                                                                                                                                                    |
|-------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Winter 2025 | CSE 582: Ethics in Artificial Intelligence, University of Washington<br>Instructor: Yulia Tsvetkov                                                                                                                                                                 |
| Spring 2024 | CSE 582: Ethics in Artificial Intelligence, University of Washington<br>Instructor: Yulia Tsvetkov<br>Redesigned the social bias lecture, updated readings each year, introduced role-based paper discussions, held office hours, and advised 20+ course projects. |

### *Research Mentees*

|                   |                                                                                   |
|-------------------|-----------------------------------------------------------------------------------|
| Christina Hahn    | Multi-LLM alignment via Stackelberg games [U4]; currently M.S. student, UW        |
| Ziyuan Yang       | Safety in model collaboration [C48, U3]; currently M.S. student, UW               |
| Zongwan Cao       | Model selection for LLM collaboration [U2]; currently M.S. student, UW            |
| Yizhuo Zhang      | Generalization of LLM graph reasoning [C45, C30] → Research Scientist, Glean      |
| Yuru Jiang        | Multi-LLM alignment through combat [C43] → Researcher, Max Planck Institute       |
| Varich Boonsanong | Fine-grained factual verification [C36] → Snowflake                               |
| Jihan Yao         | Aligning LLMs with wrong answers [C35] → Ph.D. student, UW                        |
| Zifeng Zhu        | Multimodal creativity in the wild [P2] → Ph.D. student, Xi'an Jiaotong University |
| Wenxuan Ding      | Structured knowledge reasoning [C24] → Ph.D. student, NYU                         |
| Yichen Wang       | Robustness of machine-generated text detectors [C25] → Ph.D. student, UChicago    |
| Yike Wang         | Knowledge conflicts in LLMs [C28] → Ph.D. student, UW                             |
| Yuyang Bai        | Knowledge generalization in LLMs [C21] → Ph.D. student, Texas A&M                 |
| Yuhan Liu         | Summarization, bot detection [C22, C12] → Ph.D. student, NYU                      |
| Heng Wang         | Graph reasoning in natural language [C16] → Ph.D. student, UIUC                   |
| Zhaoxuan Tan      | Bot detection, knowledge graphs [C18, C10] → Ph.D. student, Notre Dame            |
| Zhenyu Lei        | Bot detection with text-graph interaction [C13] → Ph.D. student, UVA              |

### *Undergraduate Research Lab*

|                 |                                                                                                                                                                                                                                                                                                                                   |
|-----------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| 08/2021–06/2022 | Founder and Director, <a href="#">LUD Lab</a> , Xi'an Jiaotong University. Led a lab of 19 undergraduates on knowledge-aware NLP, graph learning, and social media analysis; the lab has since grown to 40+ members, with alumni pursuing Ph.D.s at UIUC, NYU, Notre Dame, Texas A&M, UVA, Tsinghua, Peking University, and CUHK. |
|-----------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|

## **Grants & Funding**

---

### *Fellowships*

|      |                                                     |
|------|-----------------------------------------------------|
| 2026 | NVIDIA Graduate Fellowship; \$60,000                |
| 2025 | Jane Street Graduate Research Fellowship; \$112,392 |
| 2024 | IBM PhD Fellowship; \$40,000                        |

### *Grants and Gifts*

|                 |                                                                            |
|-----------------|----------------------------------------------------------------------------|
| 09/2024–Present | DARPA Scientific Feasibility (SciFy) Program; co-wrote the proposal        |
| 07/2024–12/2024 | Google Research Gift, “Teaching LLMs to Say ‘I Don’t Know’ ”; \$30,000     |
| 03/2024–09/2024 | OpenAI Researcher Access Program; \$10,000 in OpenAI API credits           |
| 02/2024–06/2024 | Microsoft Accelerate Foundation Models Research; \$20,000 in Azure credits |

## **Service**

---

### *Area Chair*

|              |                                   |
|--------------|-----------------------------------|
| 2026–Present | Senior Program Committee, AAAI    |
| 2024–Present | Action Editor, ACL Rolling Review |

### *Workshop Organizer*

|           |                                                                      |
|-----------|----------------------------------------------------------------------|
| ACL 2025  | The 3rd Workshop on Towards Knowledgeable Foundation Models (KnowFM) |
| AAAI 2025 | The 2nd Workshop on Towards Knowledgeable Foundation Models (KnowFM) |
| ACL 2024  | The 1st Workshop on Towards Knowledgeable Language Models (KnowLLM)  |

### *Session Chair*

NAACL 2025                      Machine Learning for NLP

### *Reviewer*

NLP                              ACL Rolling Review (2022–), ACL (2023–), EMNLP (2023–), COLM (2024–), COLING (2025–)

Machine Learning            NeurIPS (2022–), NeurIPS Datasets and Benchmarks Track (2022–), ICML (2022–), ICLR (2023–), Learning on Graphs Conference (2022–), IJCAI (2024–), AAAI (2025–), AISTATS (2025–)

Vision                          ECCV (2022–), CVPR (2023–), ICCV (2023–)

Web and Social                CSCW (2022), ICWSM (2023–), The Web Conference (2024–)

Journals                        IEEE Transactions on Knowledge and Data Engineering (2025–), IEEE Transactions on Neural Networks and Learning Systems (2024–), ACM Transactions on the Web (2024–), Social Network Analysis and Mining (2021)

### *Within the University*

2025–2026                      Organizer, UW NLP Seminar

2022–2025                      Mentor, Pre-Application Mentorship Service (PAMS), Paul G. Allen School

2022–2024                      Student Reviewer, Ph.D. Admissions, Paul G. Allen School

## **Selected Media Coverage**

---

05/2026                      *Allen School News*. “[Allen School Ph.D. student Shangbin Feng, 2026 NVIDIA Graduate Fellow, aims to make AI a model of collaboration.](#)” Kristine White.

12/2025                      *Science News*. “[Chatbots spewing facts, and falsehoods, can sway voters.](#)” Sujata Gupta.

08/2025                      *GeekWire*. “[UW study: Politically persuasive AI chatbots offer potential benefits—and worrying influence.](#)” Lisa Stiffler.

08/2025                      *UW News*. “[With just a few messages, biased AI chatbots swayed people’s political views.](#)” Stefan Milne.

02/2025                      *Allen School News*. “[Allen School’s Shangbin Feng and Rock Yuren Pang earn IBM Ph.D. Fellowships for seeking to change the narrative around AI.](#)” Kristine White.

08/2024                      *UW News*. “[Large language models can help detect social media bots—but can also make the problem worse.](#)” Stefan Milne.

08/2023                      *Forbes*. “[ChatGPT has liberal bias, say researchers.](#)” Emma Woollacott.

08/2023                      *Washington Post*. “[ChatGPT leans liberal, research shows.](#)” Gerrit De Vynck.

08/2023                      *MIT Technology Review*. “[AI language models are rife with different political biases.](#)” and “[Why it’s impossible to build an unbiased AI language model.](#)” Melissa Heikkilä.

08/2023                      International coverage of [\[C15\]](#) in seven languages (German, French, Italian, Spanish, Portuguese, Japanese, Chinese), including *heise online*, *El Español*, *Poder360*, and *Yahoo! News Japan*.